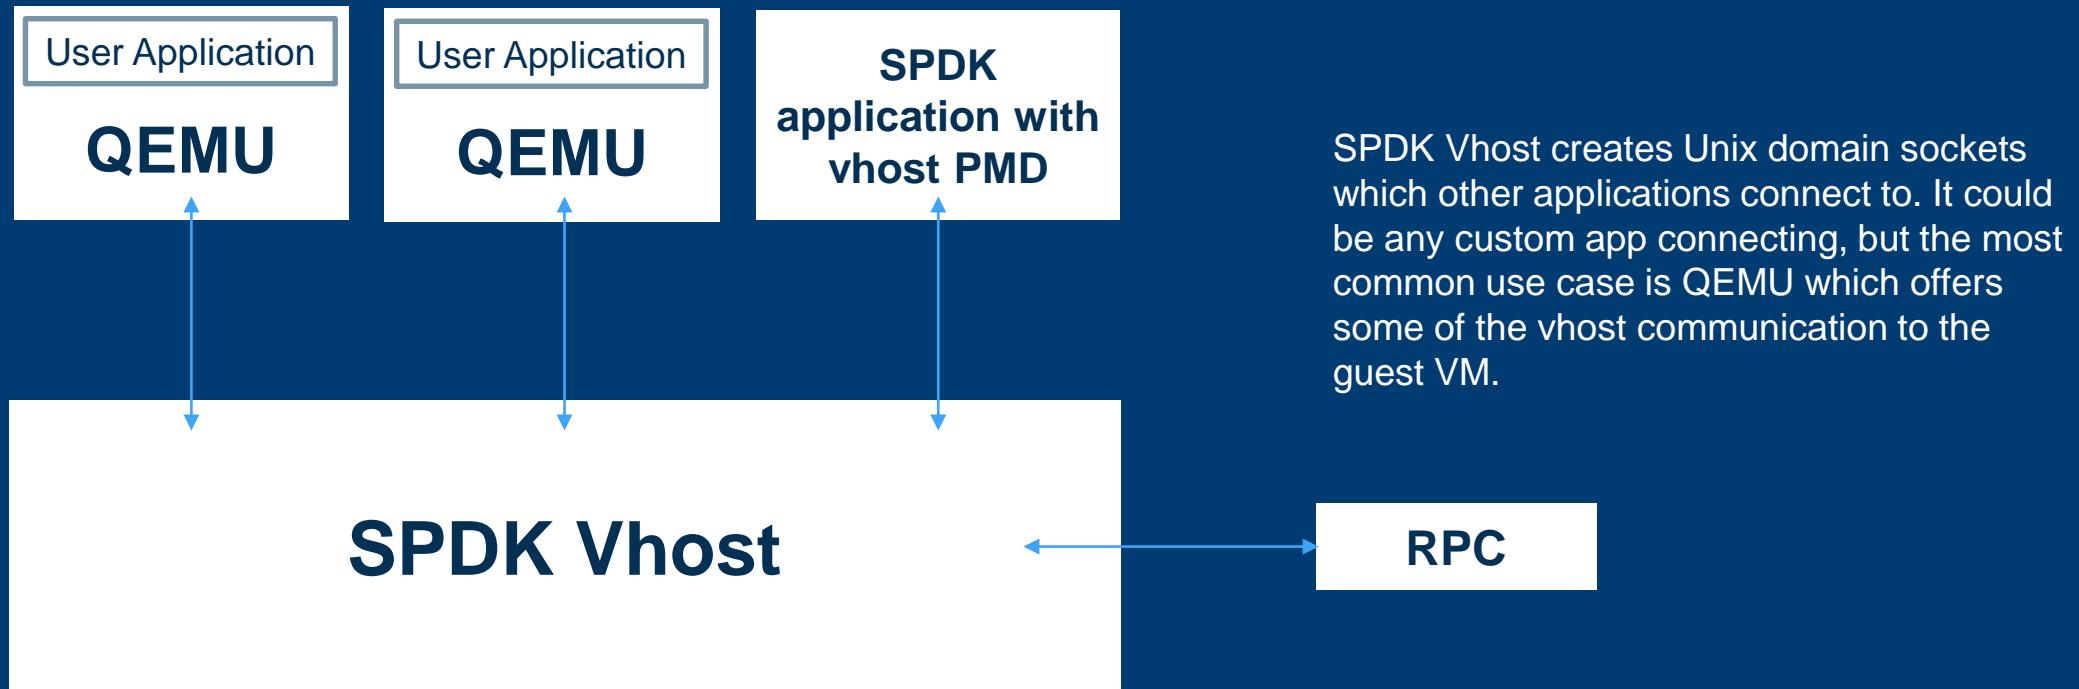
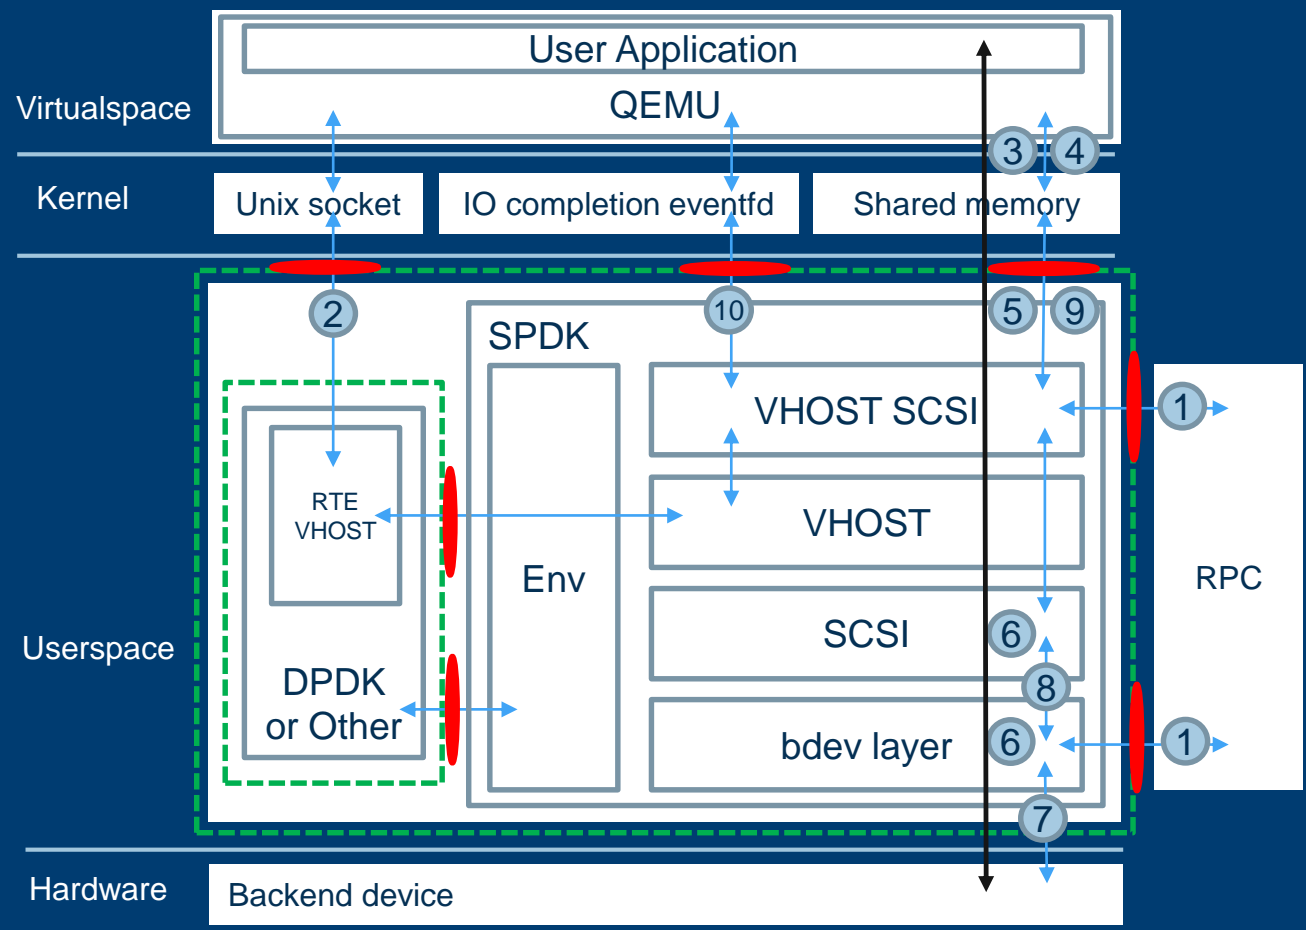


USE CASE: VHOST INTEGRATION

- User has decided to integrate SPDK vhost stack into his existing virtualization solution
- Related hardware and software components (including bdev interface) were configured correctly so they are assumed to be safe in this use case
- Users can create custom SPDK vhost applications, but this document focuses on the default one



SYSTEM DIAGRAM



- High Level Flow:
- 1) System admin starts an SPDK app and configures it through RPC
 - 2) SPDK creates a vhost Unix domain socket(s)
 - 3) System admin starts a QEMU instance and instructs it to connect to the SPDK vhost socket
 - 4) QEMU shares the entire VM's memory with SPDK
 - 5) QEMU creates a virtio-pci device to be used inside the VM
 - 6) Virtio-PCI driver inside the VM initializes the device, sets up I/O virtqueues
 - 7) QEMU, in response to device initialization, sends the relevant data over the vhost socket
 - 8) User application inside the VM sends an I/O request
 - 9) Virtio-PCI driver adds a Virtio I/O entry to a virtqueue
 - 10) SPDK detects request by continuously polling shared I/O queues
 - 11) SPDK processes the request, first in lib/scsi then in bdev layer
 - 12) Request is sent to backend device
 - 13) Callback is called from bdev layer
 - 14) Vhost updates the request status by modifying shared queues
 - 15) Vhost notifies the application about completion by writing to completion eventfd

* vhost-blk is an alternative option for vhost-scsi. If vhost-blk is used, the SCSI block can be omitted on the diagram.

ATTACK SURFACES

System Element	Compromise Type(s)	Assets exposed	Attack Method
Vhost socket	Denial of service, System integrity	RTE_VHOST, Shared memory, QEMU, SPDK app	Malformed vhost-user commands
Completion eventfd	Denial of service	Sockets, RTE_VHOST	Unexpected eventfd writes
Shared memory	Data disclosure, Denial of service, System integrity	Data, RTE_VHOST, SPDK app	Malformed IO requests
ENV/RTE VHOST interface	Data disclosure, System integrity	Env, RTE_VHOST, Data, QEMU, Sockets, Shared Memory, SPDK app	DLL Injection
RPC interface	Data disclosure, System integrity	Sockets, RTE_VHOST, SPDK app	Malformed JSON-RPC requests, Changing SPDK configuration at runtime

THREAT MATRIX

Assets Surface	Data	SPDK application	Shared memory	Sockets	Env	QEMU	RTE_VHOST
QEMU socket interface		Y		Y		Y	Y
Completion eventfd							Y
Shared memory	Y	Y	Y				Y
ENV/RTE VHOST interface	Y	Y	Y	Y	Y	Y	Y
RPC interface		Y		Y			Y

ADVERSARIES IN SCOPE

Persona	Motivation	Attacker Type	Starting Privilege Level	Skill and Potential Effort level
Malicious VM User	Wants to snoop data/disrupt users on system	Software Adversary in a VM	None	Unskilled, gives up easily
Malicious vhost client	Wants to snoop data/disrupt users on system	Vhost driver Software Adversary	None	Proficient level of skill, does not give up easily
Malicious RPC Admin	Denial Of Service	Network Adversary	None	Proficient level of skill, does not give up easily

* host system software adversary is out of scope because such adversaries have permissions to defeat all mitigations. User needs to ensure appropriate deployment policies are in place to prevent system level software adversaries

THREAT/ATTACK SURFACE MATRIX

Asset\Attack Surface	QEMU socket interface	Completion eventfd	Shared memory	ENV/ RTE_VHOST interface	RPC interface
Data availability	Y	Y	Y	Y	Y
Data confidentiality				Y	
Data integrity	Y	Y	Y	Y	Y
Shared memory resources				Y	Y
Unix sockets	Y	Y		Y	Y
App configuration file					

THREATS

Protections							
ID	Threat	Assets	Req'd	Adversary	Attack Point	Technique	Mitigation
1	Malformed vhost-user commands	data availability, sockets, shared memory	B G	Software adversary in a VM	vhost-user communication	Connect to a vhost socket and send malformed messages to stall or crash the SPDK app	SW to validate input before use
2	Malformed I/O requests	data availability, shared memory	B E	Software adversary in a VM	vhost-user communication	Connect as a client and try to setup invalid memory region to cause an error on host application	SW to validate input before use
3	Deinitialization of nonexistent virtqueues	data availability, sockets, shared memory	B G	Software adversary in a VM	vhost-user communication	Connect as a client and try to deinitialize nonexistent virtqueues to cause an error on host application	SW to validate input before use
4	Repeated reconnect	data availability	B G	Software adversary in a VM, Vhost driver software adversary	vhost-user communication	Repeatedly connect and disconnect causing SPDK to map/unmap memory regions and stall other, legitimate connections	SW to implement smart QoS and thread balancing policy

* **Protections Req'd** lists assets from system diagram

THREATS

ID	Threat	Assets	Protec tions		Attack Point	Technique	Mitigation
			Req'd	Adversary			
5	Overlapping queue addresses	data availability, sockets, shared memory	B	Software adversary in a VM	virtio data	Connect as a client and try to setup a queues with overlapping addresses to cause infinite loop or other error on host application	SW to validate input before use
6	Invalid unix socket	data availability	B G	Software adversary in a VM	vhost-user communication	Connect as a client and provide socked used for connection as ex. completion evenfd to cause infinite loop or other error on host	SW to validate input before use
7	Mutable virtio requests	data availability, data integrity, data confidentiality, sockets, shared memory	B	Vhost driver software adversary	shared memory	Modify virtio request during it being processing by host SPDK app to try to bypass error checking	SW to guarantee immutability of potentially dangerous request data such as addresses, ranges, pointers
8	Malicious RPC commands	data availability, data integrity	B G	Network adversary	RPC commands	Try to cause a race condition or other error by sending possibly conflicting RPC commands, ex hotremoving vhot controller and backend device at the same time	SW to implement privileged commands and/or monitoring system

THREATS

Protec tions							
ID	Threat	Assets	Req'd	Adversary	Attack Point	Technique	Mitigation
9	Mutable backend device pointer	data confidentiality, data availability	B	Network adversary	RPC commands	Reconfigure vhost controller to use different (available to the RPC user) backend device while it's in use, in order to get access to user data	SW to prevent mutating the controller while it is in use and/or implement privileged access to that functionality
10	Malicious vhost events	data availability, sockets, shared memory	B G	Software adversary in a VM	vhost-user communication	Connect as a client and provide socked used for connection as ex. completion evenfd to cause infinite loop or other error on host application	SW to validate input before use
11	Mutable virtio requests	data availability, data integrity, data confidentiality, sockets, shared memory	B	Vhost driver software adversary	shared memory	Modify virtio request during it being processing by host SPDK app to try to bypass error checking	SW to guarantee immutability of potentially dangerous request data such as addresses, ranges, pointers
12	WRITEs to I/O queue memory	data availability, data integrity, sockets, shared memory	B G	Vhost driver software adversary	virtio data	As virtio client, issue WRITEs to I/O queue memory to make SPDK send itself new I/Os to process and waste CPU cycles without any VM interaction	SW to implement smart QoS policy